

Databricks Assignment

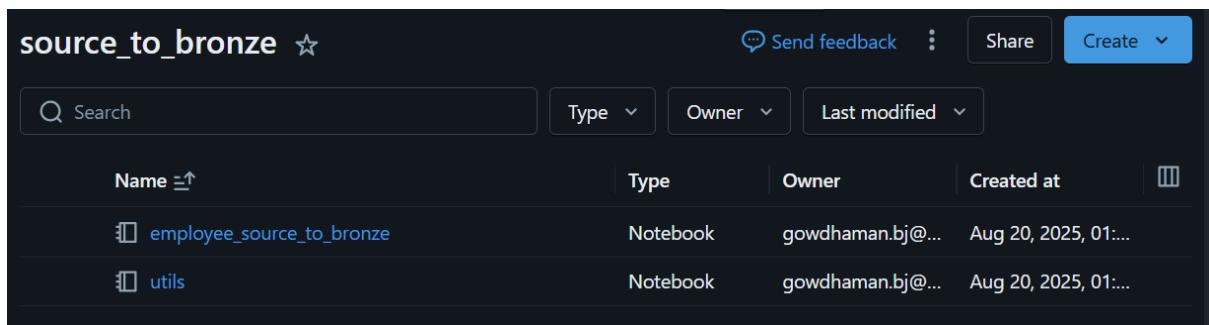
Question 1:

1. Create 3 folders as source_to_bronze, bronze_to_silver, silver_to_gold.

Name	Type	Owner	Created at
bronze_to_silver	Folder	gowdhaman.bj@...	Aug 20, 2025, 01:...
silver_to_gold	Folder	gowdhaman.bj@...	Aug 20, 2025, 01:...
source_to_bronze	Folder	gowdhaman.bj@...	Aug 20, 2025, 01:...

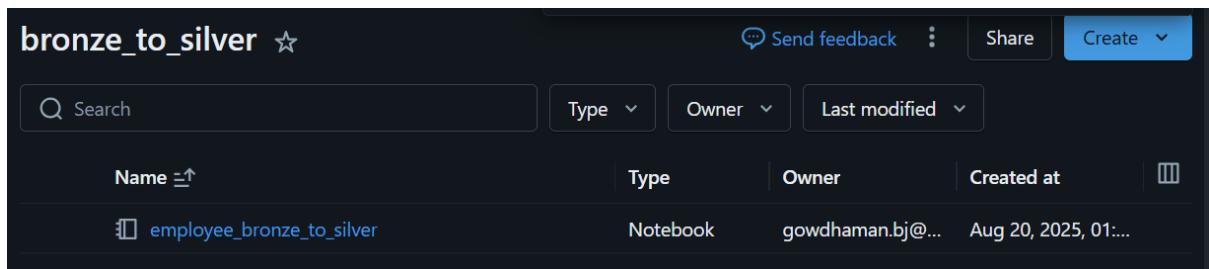
2. Create 4 notebooks in this respective order.

2 Notebooks named in source_to_bronze as utils (add all common functions in this notebook) and employee_source_to_bronze (driver notebook)



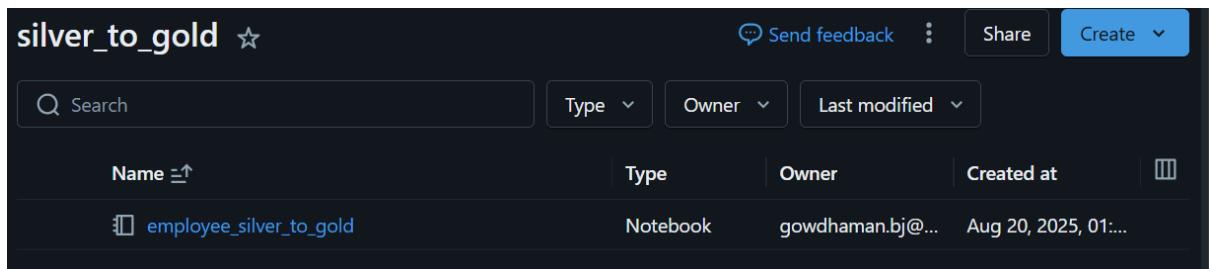
Name	Type	Owner	Created at
employee_source_to_bronze	Notebook	gowdhaman.bj@...	Aug 20, 2025, 01:...
utils	Notebook	gowdhaman.bj@...	Aug 20, 2025, 01:...

1 Notebook in bronze to silver as employee_bronze_to_silver



Name	Type	Owner	Created at
employee_bronze_to_silver	Notebook	gowdhaman.bj@...	Aug 20, 2025, 01:...

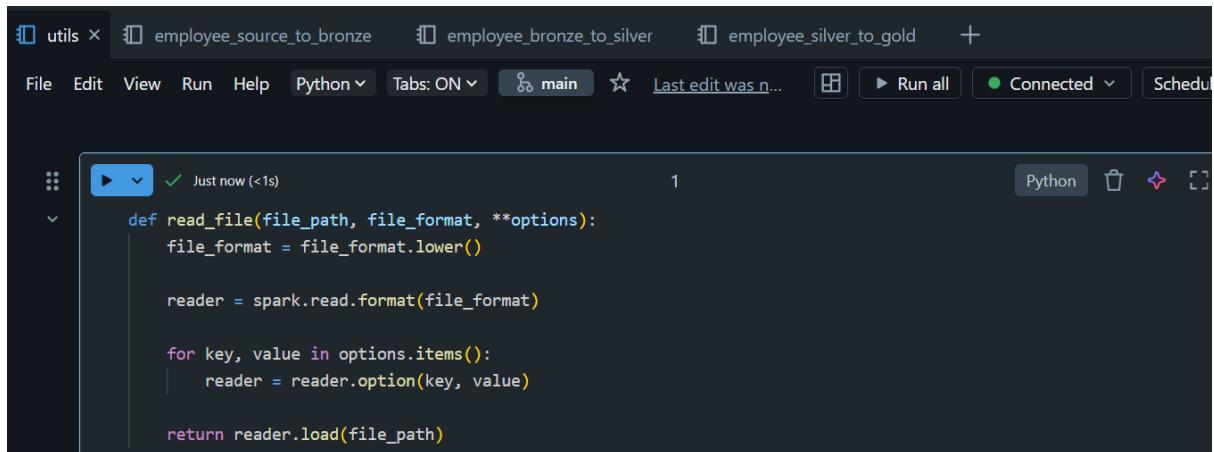
1 Notebook in silver to gold as employee_silver_to_gold



Name	Type	Owner	Created at
employee_silver_to_gold	Notebook	gowdhaman.bj@...	Aug 20, 2025, 01:...

3. Read the 3 datasets as Dataframe in employee_source_to_bronze, call utils notebook in this notebook, and write to a location in DBFS, as

/source_to_bronze/file_name.csv (employee, department_df, country_df) as CSV format.



utils x employee_source_to_bronze employee_bronze_to_silver employee_silver_to_gold +

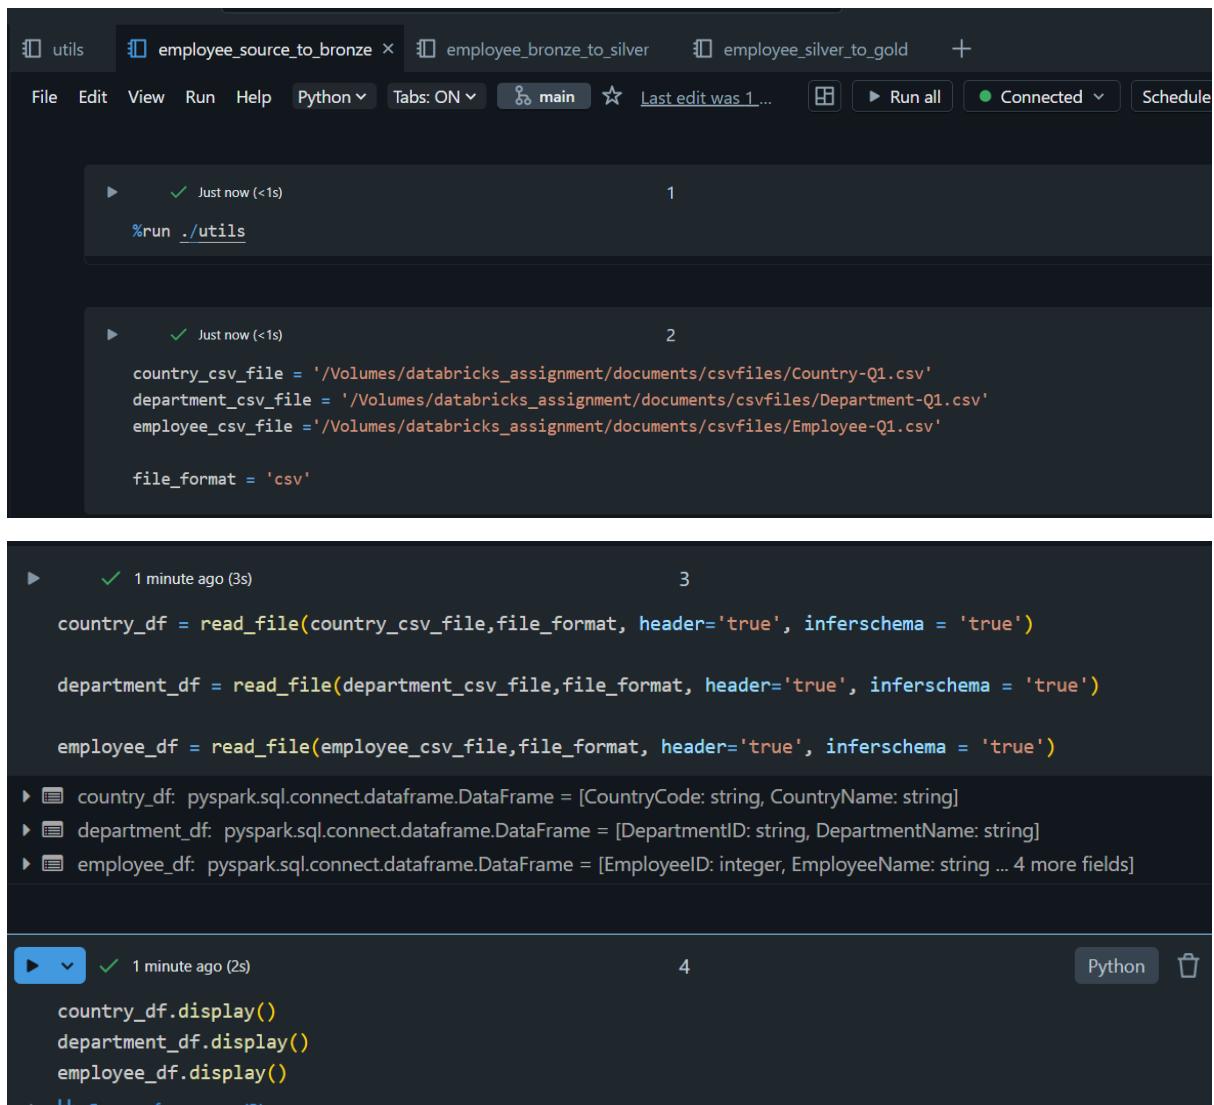
File Edit View Run Help Python Tabs: ON main Last edit was n... Run all Connected Schedule

```
▶ Just now (<1s) 1
def read_file(file_path, file_format, **options):
    file_format = file_format.lower()

    reader = spark.read.format(file_format)

    for key, value in options.items():
        reader = reader.option(key, value)

    return reader.load(file_path)
```



utils x employee_source_to_bronze x employee_bronze_to_silver employee_silver_to_gold +

File Edit View Run Help Python Tabs: ON main Last edit was 1... Run all Connected Schedule

```
▶ Just now (<1s) 1
%run ./utils
```

```
▶ Just now (<1s) 2
country_csv_file = '/Volumes/databricks_assignment/documents/csvfiles/Country-Q1.csv'
department_csv_file = '/Volumes/databricks_assignment/documents/csvfiles/Department-Q1.csv'
employee_csv_file = '/Volumes/databricks_assignment/documents/csvfiles/Employee-Q1.csv'

file_format = 'csv'
```

```
▶ 1 minute ago (3s) 3
country_df = read_file(country_csv_file, file_format, header='true', inferSchema = 'true')

department_df = read_file(department_csv_file, file_format, header='true', inferSchema = 'true')

employee_df = read_file(employee_csv_file, file_format, header='true', inferSchema = 'true')

▶ country_df: pyspark.sql.connect.DataFrame = [CountryCode: string, CountryName: string]
▶ department_df: pyspark.sql.connect.DataFrame = [DepartmentID: string, DepartmentName: string]
▶ employee_df: pyspark.sql.connect.DataFrame = [EmployeeID: integer, EmployeeName: string ... 4 more fields]
```

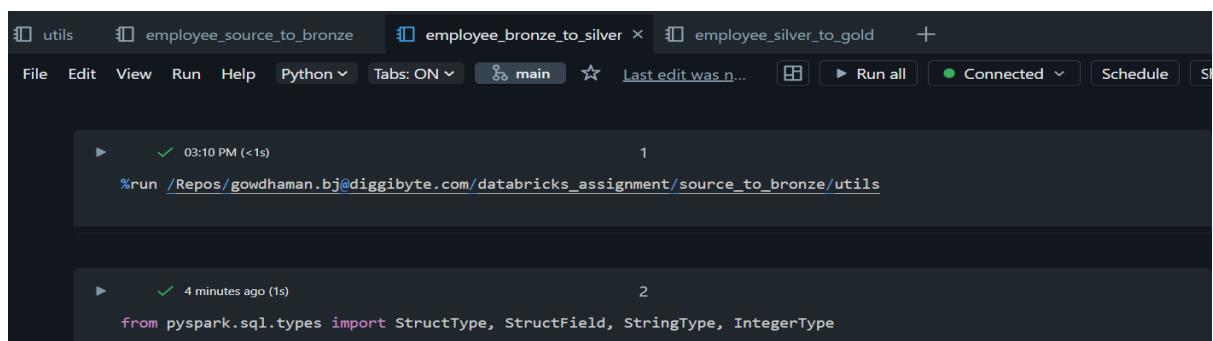
```
▶ 1 minute ago (2s) 4
country_df.display()
department_df.display()
employee_df.display()
↳ See performance (2)
```

	Δ^B_C DepartmentID	Δ^B_C DepartmentName
1	D101	Sales
2	D102	Marketing
3	D103	Finance
4	D104	Support
5	D105	HR

	Δ^B_C CountryCode	Δ^B_C CountryName
1	CN	China
2	IN	India
3	SA	South Africa
4	JA	Japan
5	MY	Malaysia
6	MA	Morocco

	Δ^2_3 EmployeeID	Δ^B_C EmployeeName	Δ^B_C Department	Δ^B_C Country	Δ^2_3 Salary	Δ^2_3 Age
1	1	James	D101	IN	9000	25
2	2	Michel	D102	SA	8000	26
3	3	James son	D101	IN	10000	35
4	4	Robert	D103	MY	11000	34
5	5	Scott	D104	MA	6000	36
6	6	Gen	D105	JA	21345	24
7	7	John	D102	MY	87654	40
8	8	Maria	D105	SA	38144	38
9	9	Soffy	D103	IN	23456	29
10	10	Amy	D103	CN	21345	24

4. In **employee_bronze_to_silver**, call utils notebook in this notebook.
 Read the file located in DBFS location source_to_bronze with as data frame different read methods using custom schema.
 (dbfs access is not their)



```
%run /Repos/gowdhaman.bj@diggibyte.com/databricks_assignment/source_to_bronze/utils
```

```
from pyspark.sql.types import StructType, StructField, StringType, IntegerType
```

▶ ✓ 3 minutes ago (<1s) 3

```
dept_schema = StructType([
    StructField("DepartmentID", IntegerType(), True),
    StructField("DepartmentName", StringType(), True)
])

employee_schema = StructType([
    StructField("EmployeeID", IntegerType(), True),
    StructField("EmployeeName", StringType(), True),
    StructField("Department", StringType(), True),
    StructField("Country", StringType(), True),
    StructField("Salary", IntegerType(), True),
    StructField("Age", IntegerType(), True)
])

country_schema = StructType([
    StructField("CountryCode", StringType(), True),
    StructField("CountryName", StringType(), True)
])
```

▶ ✓ 2 minutes ago (<1s) 4

```
country_csv_file = '/Volumes/databricks_assignment/documents/csvfiles/Country-Q1.csv'
department_csv_file = '/Volumes/databricks_assignment/documents/csvfiles/Department-Q1.csv'
employee_csv_file = '/Volumes/databricks_assignment/documents/csvfiles/Employee-Q1.csv'

file_format = 'csv'
```

▶ ✓ 1 minute ago (7s) 5

```
read_dept = read_file(department_csv_file,file_format,schema = dept_schema,header=True)
read_country = read_file(country_csv_file,file_format,schema=country_schema,header=True)
read_employee = read_file(employee_csv_file,file_format,schema=employee_schema,header=True)
display(read_dept)
display(read_country)
display(read_employee)
```

5. convert the Camel case of the columns to the snake case using UDF.

convert the Camel case of the columns to the snake case using UDF.

▶ ✓ 03:22 PM (<1s)

7

```
import re
```

▶ ✓ 5 minutes ago (<1s)

8

```
def camel_to_snake_case(camel_str):
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', camel_str)
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()

def rename_columns_to_snake_case(df):
    new_cols = [camel_to_snake_case(col) for col in df.columns]
    return df.toDF(*new_cols)
```

▶ ✓ 03:38 PM (4s)

9

```
read_dept = rename_columns_to_snake_case(read_dept)
read_country = rename_columns_to_snake_case(read_country)
read_employee = rename_columns_to_snake_case(read_employee)

# Show results
display(read_dept)
display(read_country)
display(read_employee)
```

▶ See performance (3)

▶ read_dept: pyspark.sql.connect.dataframe.DataFrame = [department_id: string
▶ read_country: pyspark.sql.connect.dataframe.DataFrame = [country_code: strin
▶ read_employee: pyspark.sql.connect.dataframe.DataFrame = [employee_id: str

Table ▾ +

	A ^B c department_id	A ^B c department_name
1	D101	Sales
2	D102	Marketing

6. Add the **load_date** column with the current date.

The primary key is EmployeeID, the Database name is Employee_info, Table name is dim_employee.

write the DF as a delta table to the location /silver/db_name/table_name.

Add the load_date column with the current date.

Just now (2s) 11

```
employee_load_date = read_employee.withColumn('load_date', current_date())
display(employee_load_date)
```

See performance (1) Optimize

employee_load_date: pyspark.sql.connect.DataFrame = [employee_id: string, employee_name: string ... 5 more fields]

Table +

employee_id	employee_name	department	country	salary	age	load_date
1	James	D101	IN	9000	25	2025-08-20
2	Michel	D102	SA	8000	26	2025-08-20
3	James son	D101	IN	10000	35	2025-08-20
4	Robert	D103	MY	11000	34	2025-08-20
5	Scott	D104	MA	6000	36	2025-08-20

The primary key is EmployeeID, the Database name is Employee_info, Table name is dim_employee.

write the DF as a delta table to the location /silver/db_name/table_name.

3 minutes ago (8s) 13

```
employee_load_date.write.format('delta').mode('overwrite').option("overwriteSchema",
"true").saveAsTable('databricks_assignment.employee_info.dim_employee')
```

7. In gold notebook employee_silver_to_gold, call utils notebook in this notebook
Read the table stored in a silver layer as DataFrame and select the columns based on
the following requirements.

```

# Read Silver table
emp_tab_df = spark.read.table("databricks_assignment.employe_info.dim_employee")
display(emp_tab_df)

```

Optimize

	Department	Country	Salary	Age	Load Date
1	D101	IN	9000	25	2025-08-20
2	D102	SA	8000	26	2025-08-20

8. Requirements:

- Find the salary of each department in descending order.

```

df_join_salary = emp_tab_df.join(department_df, emp_tab_df["Department"] == department_df["DepartmentID"], "inner")

df_join_salary.groupBy("DepartmentName").agg(
    sum("Salary").alias("Total Salary")
).orderBy("Total Salary", ascending=False).display()

```

DepartmentName	Total Salary
Marketing	95654
HR	59489
Finance	55801
Sales	19000
Support	6000

Find the number of employees in each department located in each country.

Find the number of employees in each department located in each country.

```
▶ ✓ 06:34 PM (2s) 11
df_country_join = df_join_salary.join(country_df, emp_tab_df["Country"] == country_df["CountryCode"], "inner")
▶ df_country_join: pyspark.sql.connect.dataframe.DataFrame = [employee_id: string, employee_name: string ... 9 more fields]
```

```
▶ ✓ 06:34 PM (2s) 12
display(df_country_join)
> See performance (1) Optimize
```

Just now (4s) 13 Python ⚙️ ⚡ ☰ ⋮

```
df_country_join.groupBy("CountryName", "DepartmentName").agg(count(col("employee_id")).alias("Employee_count")).sort("Employee_count").display()
> See performance (1) Optimize
```

Table +

	CountryName	DepartmentName	Employee_count
1	China	Finance	1
2	South Africa	HR	1
3	Morocco	Support	1
4	Malaysia	Marketing	1
5	Malaysia	Finance	1
6	South Africa	Marketing	1
7	India	Finance	1
8	Japan	HR	1
9	India	Sales	2

List the department names along with their corresponding country names

```
▶ ✓ 1 minute ago (2s) 15
df_country_join.select("CountryName", "DepartmentName").display()
> See performance (1)
```

Table +

	CountryName	DepartmentName
1	India	Sales
2	South Africa	Marketing
3	India	Sales
4	Malaysia	Finance
5	Morocco	Support
6	Japan	HR
7	Malaysia	Marketing
8	South Africa	HR
9	India	Finance
10	China	Finance

What is the average age of employees in each department?

What is the average age of employees in each department?

Just now (3s) 17 Python Optimize

```
df_country_join.groupBy("DepartmentName").agg(avg("age").alias("average_employee_age")).sort("average_employee_age").display()
```

See performance (1)

Table +

DepartmentName	average_employee_age
Finance	29
Sales	30
HR	31
Marketing	33
Support	36

Add the at_load_date column to data frames.

Add the at_load_date column to data frames.

Just now (2s) 19

```
def add_load_date(dataframe):
    return dataframe.withColumn("load_date", lit(current_date()))

employee_load_date = add_load_date(emp_tab_df)
country_df_load_date = add_load_date(country_df)
department_df_load_date = add_load_date(department_df)
```

employee_load_date: pyspark.sql.connect.DataFrame = [employee_id: string, employee_name: string ... 5 more fields]
country_df_load_date: pyspark.sql.connect.DataFrame = [CountryCode: string, CountryName: string ... 1 more field]
department_df_load_date: pyspark.sql.connect.DataFrame = [DepartmentID: string, DepartmentName: string ... 1 more field]

1 minute ago (3s) 20 Python

```
display(employee_load_date)
display(country_df_load_date)
display(department_df_load_date)
```

See performance (3)

Table +

EmployeeID	EmployeeName	Department	Country	Salary	Age	load_date
1	James	D101	IN	9000	25	2025-08-20
2	Michel	D102	SA	8000	26	2025-08-20
3	James son	D101	IN	10000	35	2025-08-20
4	Robert	D103	MY	11000	34	2025-08-20
5	Scott	D104	MA	6000	36	2025-08-20
6	Gen	D105	JA	21345	24	2025-08-20
7	John	D102	MY	87654	40	2025-08-20
8	Maria	D105	SA	38144	38	2025-08-20
9	Soffy	D103	IN	23456	29	2025-08-20

Table +

	A ^B CountryCode	A ^B CountryName	📅 load_date
	CN	China	2025-08-20
	IN	India	2025-08-20
	SA	South Africa	2025-08-20
	JA	Japan	2025-08-20
	MY	Malaysia	2025-08-20
	MA	Morocco	2025-08-20

Table +

	A ^B DepartmentID	A ^B DepartmentName	📅 load_date
1	D101	Sales	2025-08-20
2	D102	Marketing	2025-08-20
3	D103	Finance	2025-08-20
4	D104	Support	2025-08-20
5	D105	HR	2025-08-20

Write the df to dbfs location /gold/employee/table_name(fact_employee) with overwrite and replace where condition on at_load_date.

Question: 2

Api: <https://reqres.in/api/users?page=2>

```
▶    ✓ 09:07 PM (<1s)          2
      import requests
      import json

▶    ✓ 09:16 PM (<1s)          3
      api_url= "https://reqres.in/api/users"

▶    ✓ 09:19 PM (<1s)          4
      page = 2
      all_data = []

▶    ✓ 09:20 PM (1s)          5
      response = requests.get(api_url, params={"page": page})
      result = response.json()
      if result:
          data = result.get('data',[])
          all_data.extend(data)
          page += 1

▶    ✓ 09:20 PM (<1s)          6
      for user in all_data:
          print(user)

[{'id': 7, 'email': 'michael.lawson@reqres.in', 'first_name': 'Michael', 'last_name': 'Lawson', 'avatar': 'https://reqres.in/img/faces/7-image.jpg'},
 {'id': 8, 'email': 'lindsay.ferguson@reqres.in', 'first_name': 'Lindsay', 'last_name': 'Ferguson', 'avatar': 'https://reqres.in/img/faces/8-image.jpg'},
 {'id': 9, 'email': 'tobias.funke@reqres.in', 'first_name': 'Tobias', 'last_name': 'Funke', 'avatar': 'https://reqres.in/img/faces/9-image.jpg'}]
```

Read the data frame with a custom schema

Flatten the dataframe

```
▶ ✓ 09:24 PM (7m) 8
from pyspark.sql.types import StringType, IntegerType, StructField, StructType
schema = StructType([
    StructField("id", IntegerType(), True),
    StructField("email", StringType(), True),
    StructField("first_name", StringType(), True),
    StructField("last_name", StringType(), True),
    StructField("avatar", StringType(), True)
])
df = spark.createDataFrame(all_data, schema)
display(df)
```

	² ₃ id	^B _C email	^B _C first_name	^B _C last_name	^B _C avatar
1	7	michael.lawson@reqres.in	Michael	Lawson	https://reqres.in/img/faces/7-image.jpg
2	8	lindsay.ferguson@reqres.in	Lindsay	Ferguson	https://reqres.in/img/faces/8-image.jpg
3	9	tobias.funke@reqres.in	Tobias	Funke	https://reqres.in/img/faces/9-image.jpg
4	10	byron.fields@reqres.in	Byron	Fields	https://reqres.in/img/faces/10-image.j...
5	11	george.edwards@reqres.in	George	Edwards	https://reqres.in/img/faces/11-image.j...
6	12	rachel.howell@reqres.in	Rachel	Howell	https://reqres.in/img/faces/12-image.j...

Derive a new column from email as site_address with values(reqres.in)

Add load_date with the current date.

	² ₃ id	^B _C email	^B _C first_name	^B _C last_name	^B _C avatar	^B _C site_address	^B _C load_date
1	7	michael.lawson@reqres.in	Michael	Lawson	https://reqres.in/img/faces/7-image.jpg	reqres.in	2025-08-20
2	8	lindsay.ferguson@reqres.in	Lindsay	Ferguson	https://reqres.in/img/faces/8-image.jpg	reqres.in	2025-08-20
3	9	tobias.funke@reqres.in	Tobias	Funke	https://reqres.in/img/faces/9-image.jpg	reqres.in	2025-08-20
4	10	byron.fields@reqres.in	Byron	Fields	https://reqres.in/img/faces/10-image.j...	reqres.in	2025-08-20
5	11	george.edwards@reqres.in	George	Edwards	https://reqres.in/img/faces/11-image.j...	reqres.in	2025-08-20
6	12	rachel.howell@reqres.in	Rachel	Howell	https://reqres.in/img/faces/12-image.j...	reqres.in	2025-08-20

Write the data frame to location in DBFS as /db_name /table_name with Db_name as site_info and table_name as person_info with delta format and overwrite mode.

Write the data frame to location in DBFS as /db_name /table_name with Db_name as site_info and table_name as person_info with delta format and overwrite mode.

```
▶ Just now (3s) 13
df.write.format("delta").mode("overwrite").save("/Volumes/databricks_assignment/site_info/person_info")
> See performance (1)
```

Just now (3s) 14 Python ⚡ ⏪ ⏴ ⏵ ⏷

```
df2 = spark.read.format("delta").load("/Volumes/databricks_assignment/site_info/person_info")
display(df2)
> See performance (1) Optimize
```

df2: pyspark.sql.connect.DataFrame = [id: integer, email: string ... 5 more fields]

Table +

#	id	email	first_name	last_name	avatar	site_address
1	7	michael.lawson@reqres.in	Michael	Lawson	https://reqres.in/img/faces/7-image.jpg	reqres.in
2	8	lindsay.ferguson@reqres.in	Lindsay	Ferguson	https://reqres.in/img/faces/8-image.jpg	reqres.in
3	9	tobias.funke@reqres.in	Tobias	Funke	https://reqres.in/img/faces/9-image.jpg	reqres.in
4	10	byron.fields@reqres.in	Byron	Fields	https://reqres.in/img/faces/10-image.j...	reqres.in
5	11	george.edwards@reqres.in	George	Edwards	https://reqres.in/img/faces/11-image.j...	reqres.in
6	12	rachel.howell@reqres.in	Rachel	Howell	https://reqres.in/img/faces/12-image.j...	reqres.in

◀ ▶

↓ ▾ 6 rows | 2.77s runtime Refreshed now